

LIMENG GE

[Email](#) ◇ [LinkedIn](#) ◇ [GitHub](#) ◇ [Google Scholar](#) ◇ Available Jun–Sep 2027

EDUCATION

University of Chicago, M.S. in Computer Science

Sep 2026 – Expected Dec 2027

East China Normal University, B.A. in Philosophy

Sep 2022 – Jun 2026

New York University, Exchange Student

Sep 2024 – May 2025

SKILLS

Programming Languages Python, Java, C, C++, SQL

Machine Learning PyTorch, Hugging Face, SFT/RLHF/GRPO, Multimodal LLMs, RAG, Multi-GPU

Tools Git, Linux, Docker, AWS, PostgreSQL, Redis, Kubernetes, vLLM

PROFESSIONAL EXPERIENCE

Machine Learning Engineer Intern

Mar – Aug 2026

SHEIN *Shanghai, China*

- Shipped a **multimodal moderation model** to production: post-trained Qwen3-VL-4B on a 100K CoT dataset (MS-SWIFT, multi-GPU) to scan merchant uploads site-wide for prohibited items, automating **87%** of decisions
- Aligned the SFT checkpoint with **RLVR/GRPO**, redesigning accuracy/format/length rewards with asymmetric penalties for misses, lifting recall from 92% to **98%** at flat precision
- Engineered a 3-stage cascaded **image deduplication** service (MD5 → pHash/dHash → DINOv3 retrieval) with multi-GPU inference and resumable feature caching, raising the dedup rate from 7.3% to **29.6%** on 50K images
- Designed a two-layer **CoT auditing** tool for hallucination and result-reasoning inconsistency, flagging **7%** erroneous CoT in production and filtering **30%** low-quality samples, raising L2 model accuracy to **92%**

Research Intern

May – Oct 2025

FacePhys *Beijing, China*

- Built **WearNet**, a multimodal physiological signal database from 100+ participants, unifying wearable and lab-grade reference recordings with millisecond-level alignment and signal-quality labeling
- Developed a **UNet1D** model predicting continuous blood pressure waveforms rather than pointwise SBP/DBP estimates, enabling richer diagnostic signals for wearable health monitoring
- Introduced a combined loss with key-point error, trend deviation, and dirotic notch constraints, reducing SBP/DBP MAE by **30%** over a pointwise-loss baseline on the internal validation set

AI Data Engineering Intern

Mar – May 2025

Shanghai AI Laboratory *Remote*

- Defined a multi-dimensional evaluation framework covering **40+** Chinese cultural values, translating abstract value semantics into structured, machine-scorable metrics for LLM alignment
- Curated an SFT/RLHF dataset of 1K value-aligned texts and 800 preference pairs (chosen/rejected); the fine-tuned model improved alignment score by **15%** and reduced unsafe outputs by **25%** on an internal benchmark

PROJECTS

DeepTrace: Fault-Tolerant Deep Research Agent [\[GitHub\]](#)

Sep 2026

LangChain, FastAPI, SQLite, Tavily, Neo4j, Docker

- Built a durable agent runtime (Plan–Execute–Report–Verify) with a custom checkpointed state machine, leases, idempotent tool calls and multi-dimensional budget guards; recovered **20/20** hard-crashed runs across all 4 pipeline stages with **zero** duplicate side effects
- Replaced all-evidence prompting with per-section evidence routing over an append-only evidence ledger, cutting LLM tokens by **19.9%** and cumulative context exposure by **62%** with **100%** finding coverage and no increase in unsupported claims; added claim-level verification catching **99%** of injected citation errors

PUBLICATIONS

- Ge, L.** (2026). Consistent Biases in Large Language Models’ Syllogistic Reasoning. *AAAI 2026 Bridge on Logical and Symbolic Reasoning in Language Models*. [\[Paper\]](#) [\[GitHub\]](#)
- Qian, M., **Ge, L.** (2026). From Output Quality to Auditability: A Data Mining Agenda for LLM-Generated Artifacts. *ICDM 2026 BlueSky Track*. (Accepted)
- Ge, L.**, Qian, M. (2026). Am I the AI? How AI-Likeness Pressure Reshapes Human Writing Expression. *AAAI 2027 Artificial Intelligence for Social Impact Track*. (Under Review)